

Practicle:2 Linear Regression

2021BIT023
Vaibhav Lanjewar

Use case : We have to predict the salary using experience by using linear regresion

In []:

Import libraries

```
In [3]: import pandas as pd  
import numpy as np  
import matplotlib.pyplot as plt
```

```
In [4]: df=pd.read_csv("Salary_Data.csv")
```

```
In [5]: df.head(2)
```

```
Out[5]:
```

	YearsExperience	Salary
0	1.1	39343.0
1	1.3	46205.0

```
In [6]: df.tail(2)
```

```
Out[6]:
```

	YearsExperience	Salary
28	10.3	122391.0
29	10.5	121872.0

```
In [7]: df.isnull().sum()
```

```
Out[7]: YearsExperience    0  
Salary                  3  
dtype: int64
```

```
In [8]: df.notnull().sum()
```

```
Out[8]: YearsExperience    30  
Salary                  27  
dtype: int64
```

```
In [9]: df.isnull().sum()
```

```
Out[9]: YearsExperience    0
Salary                  3
dtype: int64
```

```
In [10]: df.dropna(inplace=True)
```


```
In [11]: df.corr() # correlation
```

```
Out[11]:
```

	YearsExperience	Salary
YearsExperience	1.00000	0.98131
Salary	0.98131	1.00000

```
In [13]: X=df.iloc[:, :-1].values # independent variable
y = df.iloc[:, -1].values # dependent variable
```

```
In [14]: # graph
plt.plot(X,y)
plt.show()
```



```
In [ ]:
```

```
In [19]: X
```

```
Out[19]: array([[ 1.1],
 [ 1.3],
 [ 1.5],
 [ 2. ],
 [ 2.2],
 [ 2.9],
 [ 3. ],
 [ 3.2],
 [ 3.2],
 [ 3.7],
 [ 3.9],
 [ 4. ],
 [ 4. ],
 [ 4.1],
 [ 5.3],
 [ 5.9],
 [ 6. ],
 [ 6.8],
 [ 7.1],
 [ 7.9],
 [ 8.2],
 [ 8.7],
 [ 9. ],
 [ 9.5],
 [ 9.6],
 [10.3],
 [10.5]])
```

```
In [16]: y
```

```
Out[16]: array([ 39343., 46205., 37731., 43525., 39891., 56642., 60150.,
 54445., 64445., 57189., 63218., 55794., 56957., 57081.,
 83088., 81363., 93940., 91738., 98273., 101302., 113812.,
 109431., 105582., 116969., 112635., 122391., 121872.])
```

```
In [20]: pip install -U scikit-learn
```

```
Requirement already satisfied: scikit-learn in c:\users\vlanj\anaconda3\lib\site-pack
ages (1.4.0)
Requirement already satisfied: numpy>=1.19.5 in c:\users\vlanj\anaconda3\lib\site-pac
kages (from scikit-learn) (1.24.3)
Requirement already satisfied: scipy>=1.6.0 in c:\users\vlanj\anaconda3\lib\site-pack
ages (from scikit-learn) (1.11.1)
Requirement already satisfied: joblib>=1.2.0 in c:\users\vlanj\anaconda3\lib\site-pac
kages (from scikit-learn) (1.2.0)
Requirement already satisfied: threadpoolctl>=2.0.0 in c:\users\vlanj\anaconda3\lib\s
ite-packages (from scikit-learn) (2.2.0)
Note: you may need to restart the kernel to use updated packages.
```

```
In [26]: from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=
```

```
In [25]: from sklearn.linear_model import LinearRegression
lm = LinearRegression()
lm.fit(X_train, y_train)
```

Out[25]:

▼ LinearRegression ⓘ ?

LinearRegression()

In [27]: `pred = lm.predict(X_test)`In [32]: `print("y_test",y_test)`
`print("X_test",X_test)`

```
y_test [ 37731. 112635.  83088.  91738.  56642.  55794.]
X_test [[1.5]
 [9.6]
 [5.3]
 [6.8]
 [2.9]
 [4.  ]]
```

In [33]: `pred`Out[33]: `array([41144.69206511, 117316.34008101, 76879.53928245, 90985.40002613,
 54310.16209255, 64654.45997125])`In [34]: `print(lm.intercept())`

```
-----
AttributeError                                Traceback (most recent call last)
Cell In[34], line 1
----> 1 print(lm.intercept())

AttributeError: 'LinearRegression' object has no attribute 'intercept'
```

In [35]: `#y=mx+c`
`#y=ax+c`
`print(lm.intercept_) # a`
`print(lm.coef_) #c`

```
27038.831321426056
[9403.90716246]
```


In [36]: `a=lm.intercept_`
`c=lm.coef_`In [37]: `print(a,c)`

```
27038.831321426056 [9403.90716246]
```

In [38]: `my_sal_pred=a*10+c # prediction of salary for 10 year exp`
`print(my_sal_pred)`

```
[279792.22037672]
```

In [39]: `# graph od train and test data`
`plt.scatter(X_train, y_train)`
`plt.plot(X_train, lm.predict(X_train))`Out[39]: `[<matplotlib.lines.Line2D at 0x29ef145a990>]`



In []: