

Practicle:3 Linear Regression + Error Detection

2021BIT023
Vaibhav Lanjewar

Use case : We have to predict the salary using experience by using linear regresion

Import libraries

```
In [1]: import pandas as pd  
import numpy as np  
import matplotlib.pyplot as plt
```

```
In [2]: df=pd.read_csv("Salary_Data.csv")
```

```
In [3]: df.head(2)
```

```
Out[3]:
```

| | YearsExperience | Salary |
|---|-----------------|---------|
| 0 | 1.1 | 39343.0 |
| 1 | 1.3 | 46205.0 |

```
In [4]: df.tail(2)
```

```
Out[4]:
```

| | YearsExperience | Salary |
|----|-----------------|----------|
| 28 | 10.3 | 122391.0 |
| 29 | 10.5 | 121872.0 |

```
In [5]: df.isnull().sum()
```

```
Out[5]: YearsExperience    0  
Salary                  3  
dtype: int64
```

```
In [6]: df.notnull().sum()
```

```
Out[6]: YearsExperience    30  
Salary                  27  
dtype: int64
```

```
In [7]: df.isnull().sum()
```

```
Out[7]: YearsExperience    0
Salary                  3
dtype: int64
```

```
In [8]: df.dropna(inplace=True)
```

```
In [9]: df.corr() # correlation
```

```
Out[9]:
```


| | YearsExperience | Salary |
|-----------------|-----------------|---------|
| YearsExperience | 1.00000 | 0.98131 |
| Salary | 0.98131 | 1.00000 |

```
In [10]: X=df.iloc[:, :-1].values # independent variable
y = df.iloc[:, -1].values # dependent variable
```

```
In [34]: X
```

```
Out[34]: array([[ 1.1],
 [ 1.3],
 [ 1.5],
 [ 2. ],
 [ 2.2],
 [ 2.9],
 [ 3. ],
 [ 3.2],
 [ 3.2],
 [ 3.7],
 [ 3.9],
 [ 4. ],
 [ 4. ],
 [ 4.1],
 [ 5.3],
 [ 5.9],
 [ 6. ],
 [ 6.8],
 [ 7.1],
 [ 7.9],
 [ 8.2],
 [ 8.7],
 [ 9. ],
 [ 9.5],
 [ 9.6],
 [10.3],
 [10.5]])
```

```
In [11]: # graph
plt.plot(X,y)
plt.show()
```



```
In [13]: from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=
```

```
In [14]: from sklearn.linear_model import LinearRegression
lm = LinearRegression()
lm.fit(X_train, y_train)
```

```
Out[14]: ▼ LinearRegression ⓘ ?
LinearRegression()
```

```
In [15]: pred = lm.predict(X_test)
```

```
In [16]: print("y_test",y_test)
print("X_test",X_test)

y_test [ 37731. 112635.  83088.  91738.  56642.  55794.]
X_test [[1.5]
 [9.6]
 [5.3]
 [6.8]
 [2.9]
 [4.  ]]
```

```
In [17]: pred
```

```
Out[17]: array([ 41144.69206511, 117316.34008101,  76879.53928245,  90985.40002613,
 54310.16209255,  64654.45997125])
```

$$y=mx+c$$

$$y=ax+c$$

In [18]: `print(lm.intercept())`

```
-----
AttributeError                                Traceback (most recent call last)
Cell In[18], line 1
----> 1 print(lm.intercept())

AttributeError: 'LinearRegression' object has no attribute 'intercept'
```

In [19]: `a=lm.intercept_
c=lm.coef_`

In [20]: `print(a,c)`


27038.831321426056 [9403.90716246]

In [21]: `my_sal_pred=a*10+c # prediction of salary for 10 year exp
print(my_sal_pred)`

[279792.22037672]

In [22]: `# graph od train and test data
plt.scatter(X_train, y_train)
plt.plot(X_train, lm.predict(X_train))`

Out[22]: [`<matplotlib.lines.Line2D at 0x17e1718e010>`]



Model Evaluation Metrix

```
In [23]: from sklearn import metrics
```

```
In [24]: print('Mean Absolute Error is : ',metrics.mean_absolute_error(y_test,pred))
```

Mean Absolute Error is : 4374.731786040949

```
In [25]: print('Mean Squared Error is : ',metrics.mean_squared_error(y_test,pred))
```

Mean Squared Error is : 26104141.43339284

```
In [26]: print('Root Mean Squared Error is: ',np.sqrt(metrics.mean_squared_error(y_test,pred)))
```

Root Mean Squared Error is: 5109.2212159381825

```
In [27]: #root Mean Squared Error (RMSE)
```

```
In [28]: train_score_lm = lm.score(X_train, y_train)
test_score_lm = lm.score(X_test, y_test)
```

```
print("Train score: ", train_score_lm)
print("Test score : ",test_score_lm)
```

Train score: 0.9633907320629322

Test score : 0.9591199103412812

```
In [ ]: from sklearn.metrics import r2_score
print(" Root mean Squared error is:",r2_score(y_test,pred) )
r2=r2_score(y_test,pred)
```

Pracrice:4 Multiple Linear Regression

```
In [32]: from sklearn.compose import ColumnTransformer
from sklearn.preprocessing import OneHotEncoder
col_transf = ColumnTransformer(transformers=[('encoder', OneHotEncoder(), [])], remain
X = np.array(col_transf.fit_transform(X))
```

```
In [ ]:
```

```
In [33]: X
```

```
Out[33]: array([[ 1.1],
 [ 1.3],
 [ 1.5],
 [ 2. ],
 [ 2.2],
 [ 2.9],
 [ 3. ],
 [ 3.2],
 [ 3.2],
 [ 3.7],
 [ 3.9],
 [ 4. ],
 [ 4. ],
 [ 4.1],
 [ 5.3],
 [ 5.9],
 [ 6. ],
 [ 6.8],
 [ 7.1],
 [ 7.9],
 [ 8.2],
 [ 8.7],
 [ 9. ],
 [ 9.5],
 [ 9.6],
 [10.3],
 [10.5]])
```

```
In [35]: from sklearn.model_selection import train_test_split
Xm_train, Xm_test, ym_train, ym_test = train_test_split(X, y, test_size = 0.3, random_
```

```
In [36]: from sklearn.linear_model import LinearRegression
lm1 = LinearRegression()
lm1.fit(Xm_train, ym_train)
```

```
Out[36]: ▼ LinearRegression ⓘ ?
LinearRegression()
```

```
In [37]: ym = lm1.predict(Xm_test)
np.set_printoptions(precision=2)
print(np.concatenate((ym.reshape(len(ym),1), ym_test.reshape(len(ym_test),1)),1))

[[ 64179.34  63218. ]
 [113452.36 105582. ]
 [127944.42 121872. ]
 [ 84468.23  93940. ]
 [ 66111.61  57081. ]
 [119249.19 112635. ]
 [ 57416.37  54445. ]
 [ 62247.06  57189. ]
 [118283.05 116969. ]]
```

```
In [ ]:
```